

02-05: Vector Operations and NA

1 - Purpose

- performing mathematics on vectors
- dealing with missing value in a vector
- R shortcuts for vector mathematics

2 - Concepts

3 - Vector math

For this lesson, we are going to use the data from the file **LansingWeather2.csv**. First we need to open the CSV file and save the data in the file to a data frame, which we will call ***weatherData***.

```
1 {  
2   rm(list=ls()); options(show.error.locations = TRUE);  
3  
4   # read data from LansingWeather2.csv and save to the variable weatherData...  
5   weatherData = read.csv("data/LansingWeather2.csv");  
6 }
```

The data frame, ***weatherData***, has the date, high temperature, low temperature, and precipitation for 14 days (March 27-April 9, 2017) in Lansing, MI.

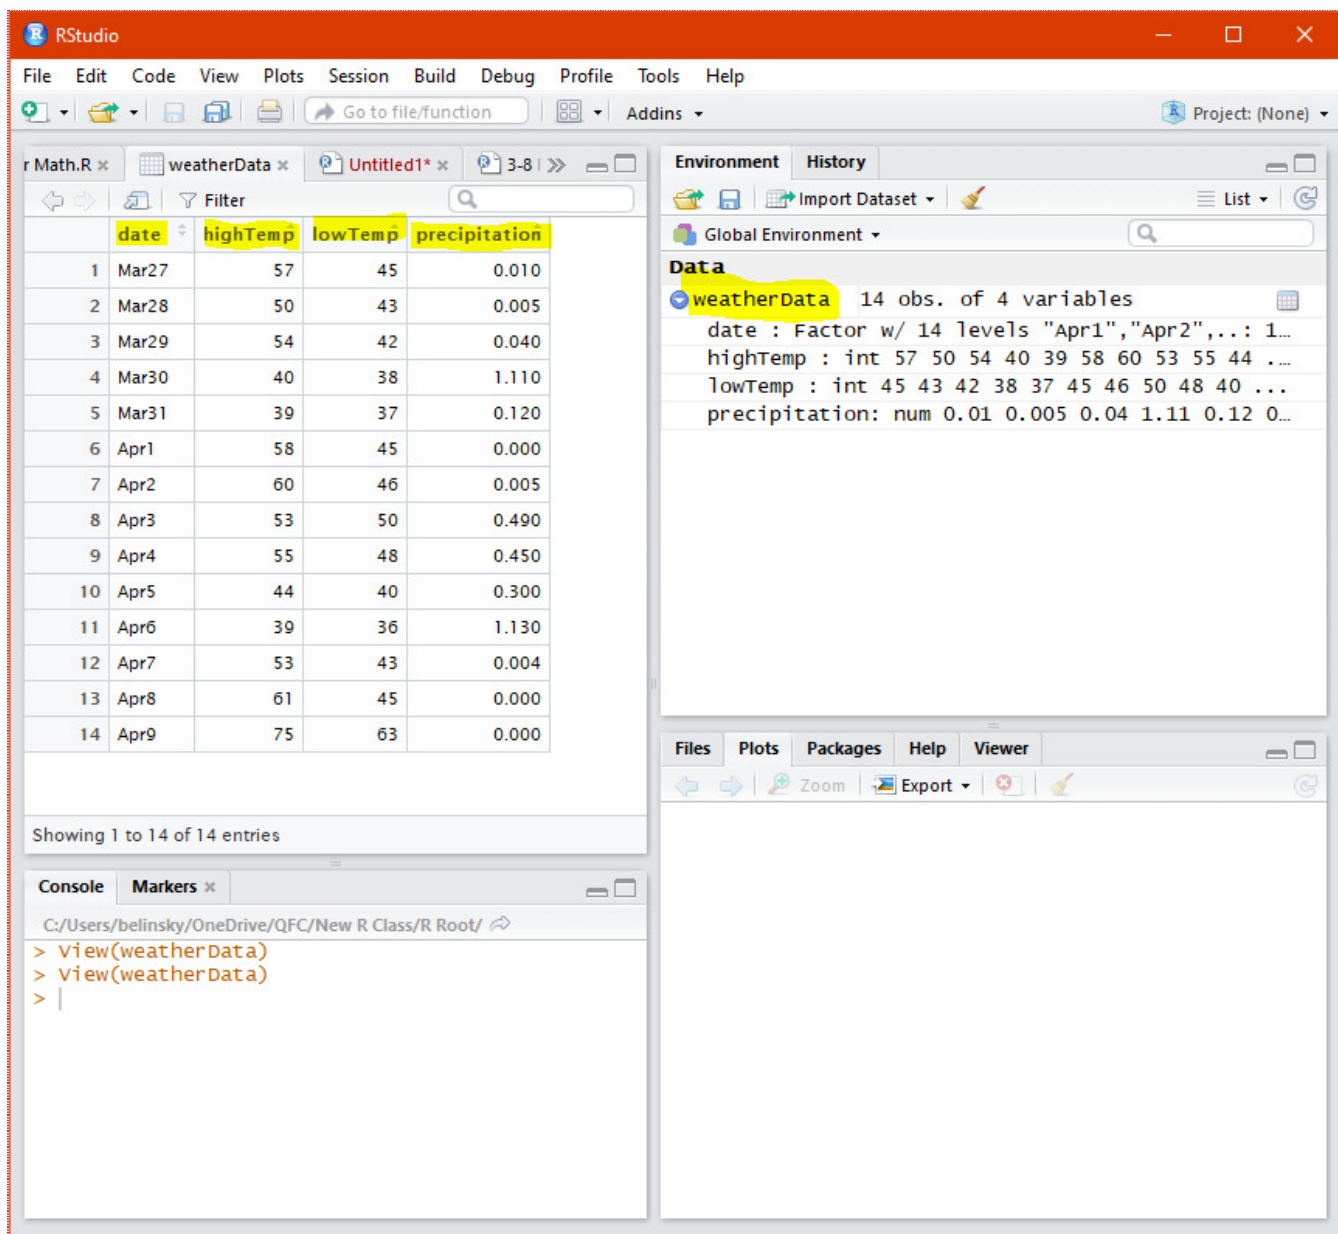


Fig 1: Looking at the value in the **weatherData** vector in the Main Window.

3.1 - Finding change in temperature

We want to find the change in temperature for each day -- in other words, the high temperature minus the low temperature. Basically this means subtracting each day's low temperature from the high temperature and saving that value to a vector.

First we will extract the low and high temperature values from **weatherData**. There are two ways we can do this:

1) by column number

```
1 # get all values from the 3rd column in weatherData
2 highTemp = weatherData[, 3];
```

2) or by column name

```
1 # get all values from the column named "lowTemp" in weatherData
```

```
2 | lowTemp = weatherData[, "lowTemp"];
```

Note: the row number is left blank to indicate we are getting values *from all rows*.

We will use a **for()** to iterate through each value (i.e., the highTemp - lowTemp for each day) and solve for the temperature difference. To do this we also need a vector that will hold the change in temperature values, which we will call **changeInTemp**.

```
1 | changeInTemp = c(); #declared a vector
```

changeInTemp acts as a state variable because:

- 1) **changeInTemp** gets initialized before the **for()**
- 2) **changeInTemp** gets populated during the iterations of the **for()**
- 3) **changeInTemp** final state is the full set of temperature changes

3.2 - Subtracting values from two different vectors

We are going to use a **for()** loop that iterates through each value in **highTemp** and **lowTemp**, subtracts the values, and saves the answer to **changeInTemp**. We need to know the length of the vectors to do this. In this case, we know there are **14** values, but we will use **length()** to get this value so that it works for vectors of any size.

```
1 | vectorLength = length(lowTemp);
```

Note: since the length of all columns in a data frame are, by definition, the same, we only need to get the length of one vector.

We use **vectorLength** to create a sequence that the **for()** iterates through. The **for()** iterates through the sequence **1:vectorLength**, or **1:14**, and assigns the value of the **14 highTemp - lowTemp** operations to the **14 changeInTemp** values.

```
1 | for(i in 1:vectorLength)
2 | {
3 |   changeInTemp[i] = highTemp[i] - lowTemp[i];
4 | }
```

Let's put all this code together:

```
1 | {
2 |   rm(list=ls()); options(show.error.locations = TRUE);
3 |
4 |   # read data from Lansingweather2.csv and save to variable weatherData
5 |   weatherData = read.csv("data/Lansingweather2.csv");
6 |   # get all values from the 3rd column in weatherData
7 |   highTemp = weatherData[, 2];
8 |   # get all values from the column named "lowTemp" in weatherData
9 |   lowTemp = weatherData[, "lowTemp"];
10 |
11 |   changeInTemp = c(); # declare a vector
```

```

12  vectorLength = length(lowTemp); # vectorLength will be 14 (length of data)
13
14  # go through the sequence 1:14
15  for(i in 1:vectorLength)
16  {
17      # subtract lowTemp from highTemp for all 14 values and save to changeInTemp
18      changeInTemp[i] = highTemp[i] - lowTemp[i];
19  }
20 }

```

The Environment Window shows that the values in the vector **changeInTemp** are **highTemp - lowTemp** but the Environment Window only shows up to 10 values. If you want to see all the values in **changeInTemp**, you can type **changeInTemp** in the Console Window and the 14 values in **changeInTemp** will appear on the next line.

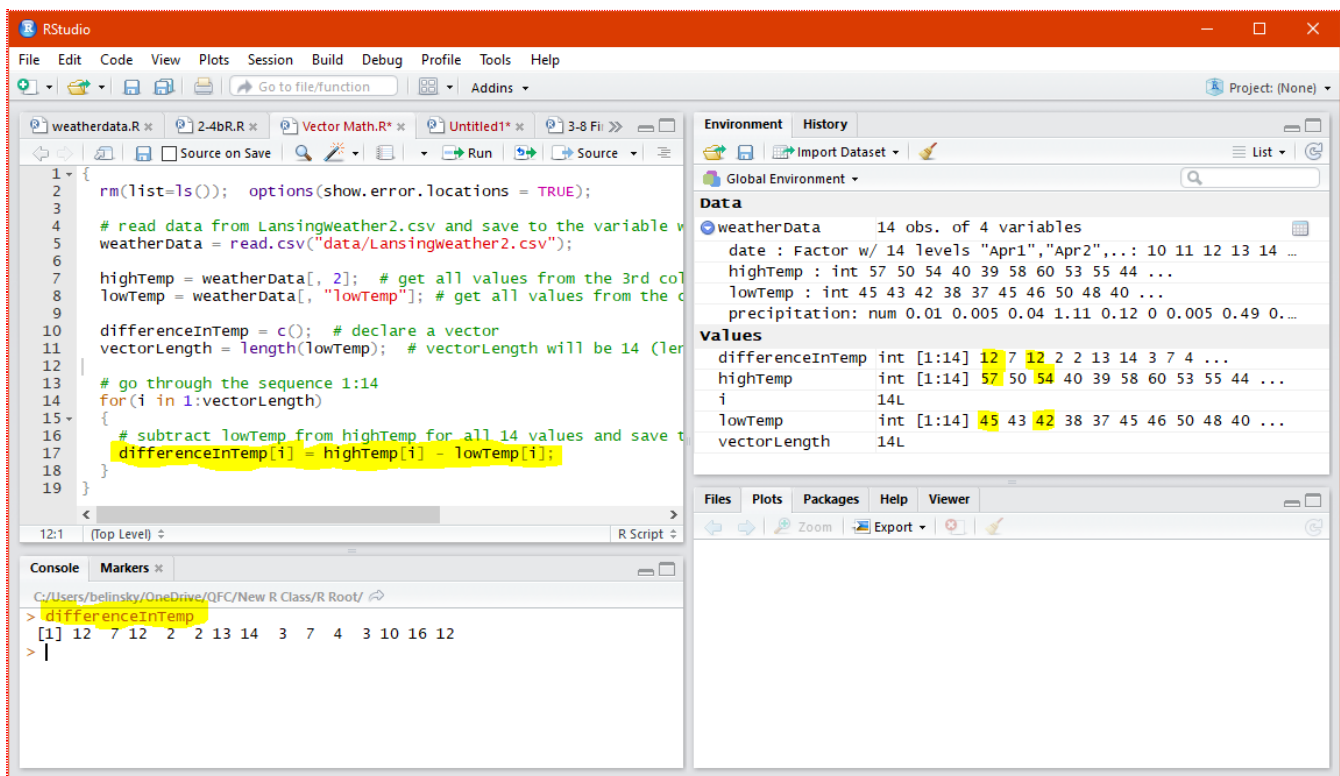


Fig 2: Using a **for()** to perform iterative mathematical operations on vectors.

4 - Dealing with missing values

In the previous example, we have an idealized situation where every day had a high and low temperature associated with it. In the real world, especially with large amounts of data, there is often missing data. In R, missing data is represented by **NA**, but the CSV file could designated missing data according to a number of different convention.

First lets create a data set with missing values and save this as **MissingTemps.csv** in the **Data** directory

```

1  date,highTemp, lowTemp, precipitation
2  Mar27,57,45,0.01

```

```

3 Mar28,50,43,0.005
4 Mar29,54, ,0.04
5 Mar30,40,38,1.11
6 Mar31,39,NA,0.12
7 Apr1,58,45,0
8 Apr2,60, ,0.005
9 Apr3,53,50,0.49
10 Apr4,55,48,0.45
11 Apr5,44,40,0.30
12 Apr6,39,36,1.13
13 Apr7,NULL,43,0.004
14 Apr8,61,45,0
15 Apr9,75,63,0

```

In the above data set, the **lowTemp** for **Mar29** and **Apr2** are left blank, the **lowTemp** for **Mar31** is given as **NA**, and the **highTemp** for **Apr7** is given as **NULL**. So, there are four missing values in **missingTemps.csv**.

4.1 - Viewing blank or NA values

We are going to open a new script and save **missingTemps.csv** data to a Data Frame called **weatherData**. But, we want to standardized the way that missing values are recorded. So, in **read.csv()** we add the parameter **na.strings**:

```

1 na.strings=c("", " ", "NULL", NULL, "NA", "na", NA, "null")

```

Essentially, **na.strings** is a vector that contains all the value that you want R to assign as **NA**. It is good to be paranoid here and think of all the possible ways in which CSV files will present these values!

```

1 {
2   rm(list=ls()); options(show.error.locations = TRUE);
3
4   weatherData = read.csv("data/missingTemps.csv",
5     na.strings=c("", " ", "NULL", NULL, "NA", "na", NA, "null"));
6 }

```

Execute the code above and double-click on **weatherData** in the **Environment Window** so that it appears in the Main Window. We see that all four missing values in **weatherData** are labelled as **NA**.

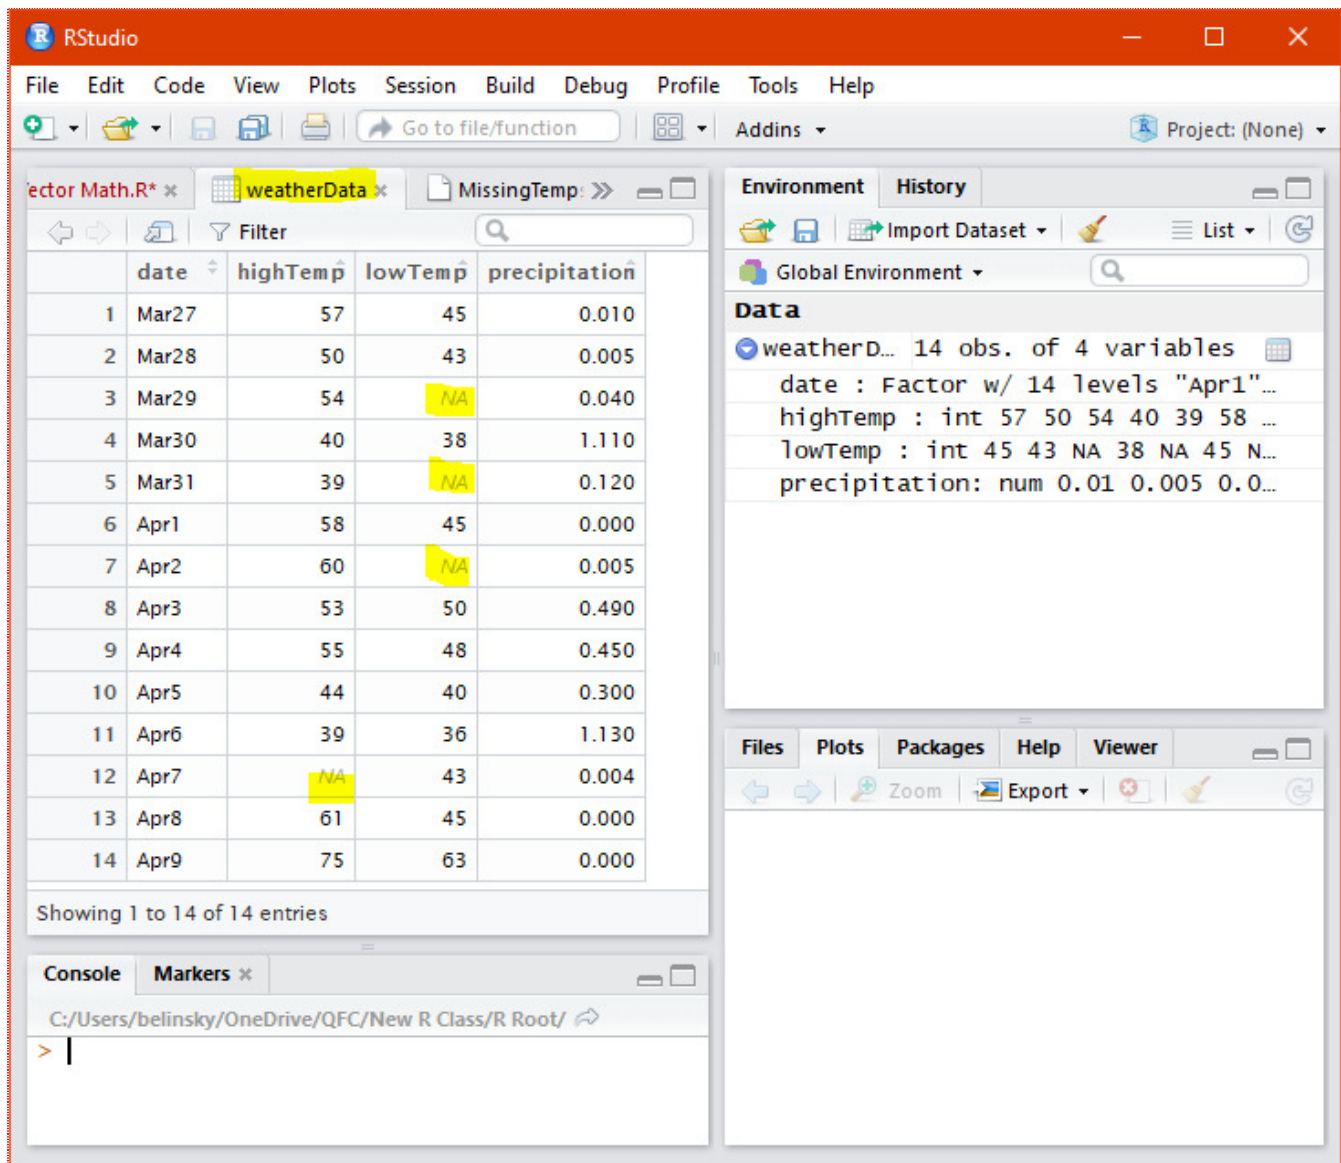


Fig 3: NA values in a data frame

4.2 - Mathematical operation on vectors with NA values

We are going to run the same change in temperature script as before except with NA values in the data frame:

```

1 {
2   rm(list=ls()); options(show.error.locations = TRUE);
3
4   # read data from missingTemps.csv and save to the variable weatherData...
5   weatherData = read.csv("data/missingTemps.csv",
6     na.strings=c("", " ", "NULL", NULL, "NA", "na", NA, "null"));
7
8   # get all values from the 3rd column in weatherData
9   highTemp = weatherData[, 2];

```



```

10 # get all values from the column named "lowTemp" in weatherData
11 lowTemp = weatherData[, "lowTemp"];
12
13 changeInTemp = c(); # declare a vector
14 vectorLength = length(lowTemp); # vectorLength will be 14 (length of data)
15
16 # go through the sequence 1:14
17 for(i in 1:vectorLength)
18 {
19     # subtract lowTemp from highTemp for all 14 values and save to changeInTemp
20     changeInTemp[i] = highTemp[i] - lowTemp[i];
21 }
22 }

```

Note: anytime there is an **NA** in a calculation (e.g., 3rd element in **differenceInTemp**), the answer is going to be **NA**. This is a special feature of **NA**, if you used any other value to represent missing data then you would get an error in the calculation.

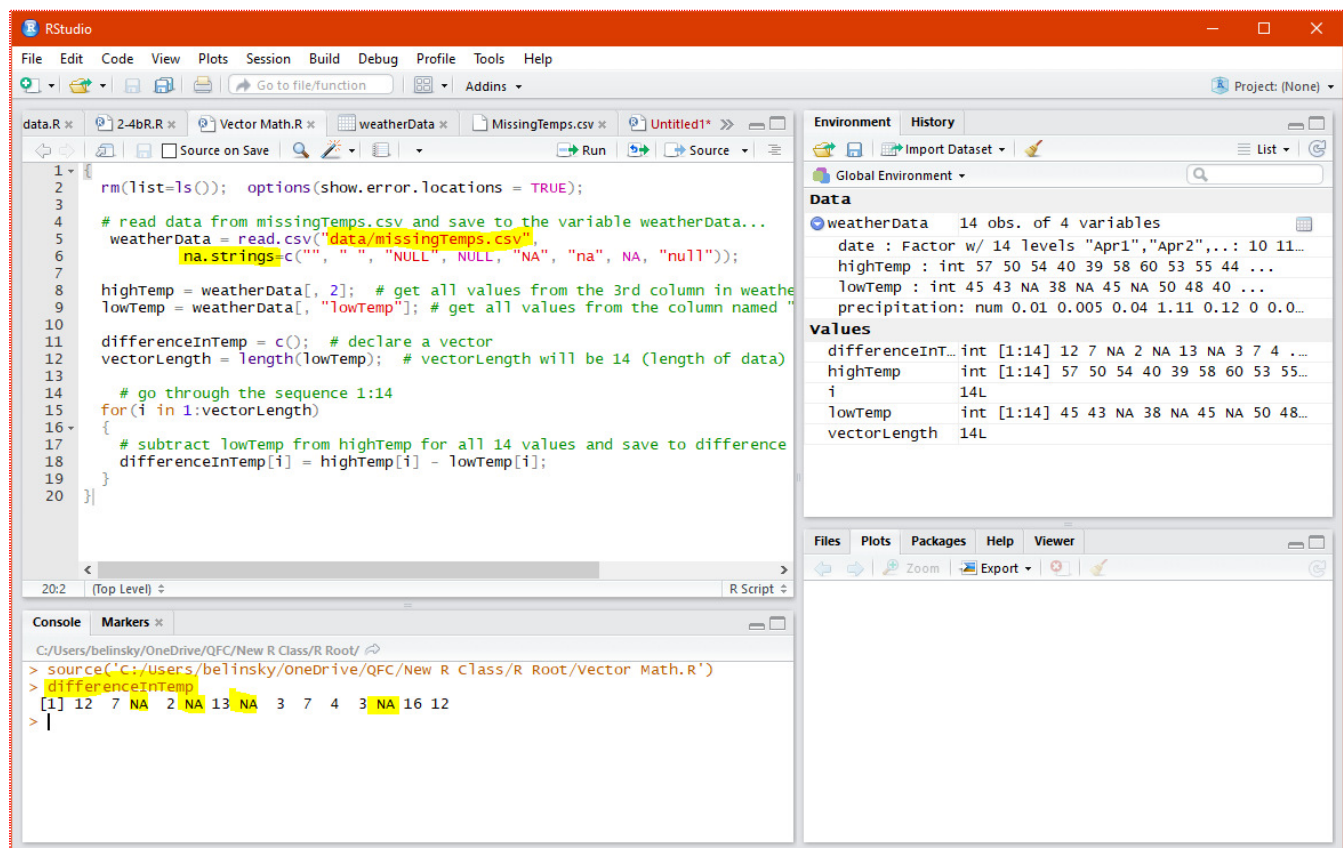


Fig 4: Doing mathematical operation with **NA** values

Extension: What counts as NA

5 - Vector Operation Shortcuts

So far we have used **for()** loops to perform iterative operations on a vector or multiple vectors. However, R

has built in functions that do all these operations using a lot less code. For example, we can perform the above change in temperature calculation without a **for()**.

```

1 {
2   rm(list=ls()); options(show.error.locations = TRUE);
3
4   weatherData = read.csv("data/missingTemps.csv",
5     na.strings=c("", " ", "NULL", NULL, "NA", "na", NA, "null"));
6
7   highTemp = weatherData[, 2];
8   lowTemp = weatherData[, "lowTemp"];
9
10  changeInTemp = highTemp - lowTemp;
11 }

```

Line 10 does all the work of iterating through the values in the vector, subtracting the values, and saving the answer to the vector **changeInTemp**. The result is the same results as the script with the **for()** loop.

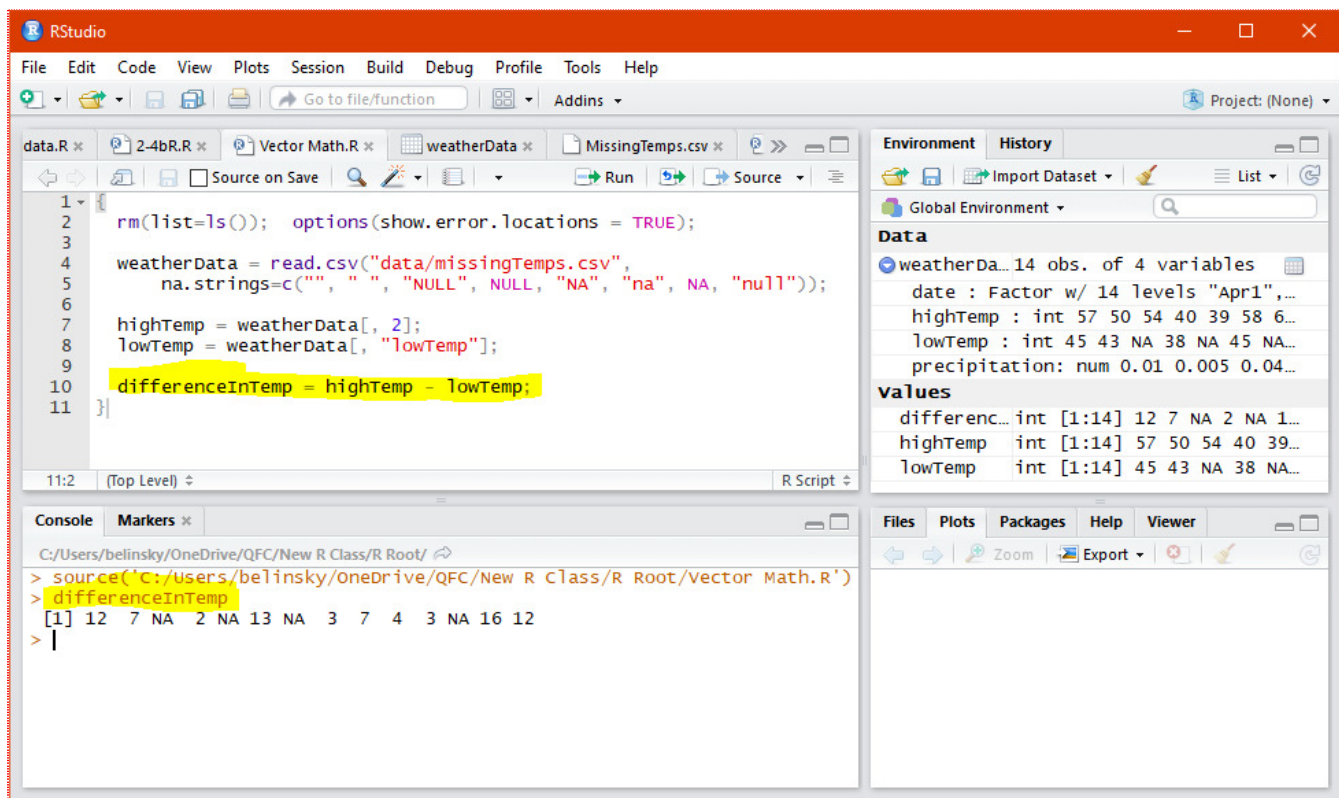


Fig 5: Subtracting a vector in R.

Obviously, this method is easier and more often used than **for()** loops for subtracting two vectors. However, this method is really a shortcut and a shortcut specific to the R language. It is important to understand what is going on in with the **for()** example because, eventually, you will find a situation where you need to use a **for()** to solve a problem.

5.1 - Many other shortcuts

R has many functions that can quickly perform the most common operations on vectors like ***sum()***, ***max()***, ***min()***, and ***mean()***.

One important parameter used in all of these functions is ***na.rm***:

na.rm = TRUE tells R to exclude the ***NA*** values from the vector before performing the operation.

na.rm = FALSE tells R to return ***NA*** if there are ***NA*** values in the vector.

The default value for ***na.rm*** is ***FALSE***.

```
1 {
2   rm(list=ls()); options(show.error.locations = TRUE);
3
4   weatherData = read.csv("data/missingTemps.csv",
5     na.strings=c("", " ", "NULL", NULL, "NA", "na", NA, "null"))
6
7   # save high temp and low temp columns to vectors
8   highTemp = weatherData[, 2];
9   lowTemp = weatherData[, "lowTemp"];
10
11  # get min and max temp (na.rm = default = FALSE)
12  minTemp = min(lowTemp);
13  maxTemp = max(highTemp);
14
15  # get min and max temp ignoring the NAs
16  minTempTake2 = min(lowTemp, na.rm=TRUE);
17  maxTempTake2 = max(highTemp, na.rm=TRUE);
18
19  # get the index of min and max values
20  minIndex = which.min(lowTemp);
21  maxIndex = which.max(highTemp);
22
23  # get simple statistics on a vector
24  sumTemp = sum(highTemp, na.rm=TRUE);
25  meanTemp = mean(highTemp, na.rm=TRUE);
26  medianTemp = median(highTemp, na.rm=TRUE);
27  stanDev = sd(highTemp, na.rm=TRUE);
28  variance = var(highTemp, na.rm=TRUE);
29
30  # get index of values that meet a condition
31  # this is more advanced and will be covered in detail later
32  whichHighGT60 = which(highTemp > 60);
33  whichLowLT43 = which(lowTemp < 43);
```

34 }

The screenshot shows the RStudio interface. The script editor on the left contains R code for reading a CSV file, extracting high and low temperature columns, and calculating various statistics. The Environment pane on the right displays the 'Global Environment' with a 'Data' section showing the 'weatherData' object (14 observations, 4 variables) and a 'values' section listing various statistical measures.

```

1 {
2   rm(list=ls()); options(show.error.locations = 1)
3
4   weatherData = read.csv("data/missingTemps.csv")
5   na.strings=c("", " ", "NULL", NULL, "NA")
6
7   # save high temp and low temp columns to vector
8   highTemp = weatherData[, 2];
9   lowTemp = weatherData[, "lowTemp"];
10
11  # get min and max temp
12  minTemp = min(lowTemp);
13  maxTemp = max(highTemp);
14
15  # get min and max temp ignoring the NAs
16  minTempTake2 = min(lowTemp, na.rm=TRUE);
17  maxTempTake2 = max(highTemp, na.rm=TRUE);
18
19  # get the index of min and max values
20  minIndex = which.min(lowTemp);
21  maxIndex = which.max(highTemp);
22
23  # get simple statistics on a vector
24  sumTemp = sum(highTemp, na.rm=TRUE);
25  meanTemp = mean(highTemp, na.rm=TRUE);
26  medianTemp = median(highTemp, na.rm=TRUE);
27  stanDev = sd(highTemp, na.rm=TRUE);
28  variance = var(highTemp, na.rm=TRUE);
29
30  # get index of values that meet a condition
31  # this is more advanced and will be covered in
32  whichHighGT60 = which(highTemp > 60);
33  whichLowLT43 = which(lowTemp < 43);
34 }

```

Environment | History

Global Environment

Data

weatherData 14 obs. of 4 variables

values

highTemp	int [1:14]	57 50 54 40 39 58 60 5...
lowTemp	int [1:14]	45 43 NA 38 NA 45 NA 5...
maxIndex		14L
maxTemp		NA_integer_
maxTempTake2		75L
meanTemp		52.6923076923077
medianTemp		54L
minIndex		11L
minTemp		NA_integer_
minTempTake2		36L
stanDev		10.3793433911192
sumTemp		685L
variance		107.730769230769
whichHighGT60	int [1:2]	13 14
whichLowLT43	int [1:3]	4 10 11

Console | Markers

C:/Users/belinsky/OneDrive/QFC/New R Class/R Root/

> |

Fig 6: Using R functions to perform common vector operations.

6 - Application

Create a vector called **changeInHighTemp** and using the data from **lansingWeather2.csv**, find the change in high temperatures from day-to-day. Add the vector to the data frame -- this means the vector must be the same size as the other columns.

So if you have four high temperatures: 40, 45, 35, 42
changeInHighTemp would be: **NA**, 5, -10, 7

The NA says that there is not enough day to give the change in temperature for the first day.

7 - Extension: What counts as NA

R applies some intelligence when it moves data from a CSV to a data frame. If there is an empty value in a column that is all numbers, then R will convert the empty value to **NA**. **NULL**, however, is not considered to be a valid data frame value and if there is a **NULL** in a column, it will be treated as the string value "NULL". R also does not consider lowercase **na** to be **NA** and this will also be treated as the string value "na".

There is now a second problem: if there is even one string in a column, all values in the column are treated as strings. This means that **40** will be seen as the string "40" and mathematical operations on strings will cause an error.