

## 03-08 Complex Conditional Subsets

### 1 - Purpose

- index a vector using multiple conditions
- index a string vector using patterns in the string values
- subset a vector using the index vector
- revisit linear regression models and scatterplots

### 2 - Concepts

### 3 - Gathering data

[The script for this lesson is located here.](#)

Like the last lesson, we are going to use the reformatted Lansing weather data frame and we are going to treat all the columns with string values as strings -- instead of factors.

```
1  lansing2016weather = read.csv(file="data/LansingNOAA2016Formatted.csv",  
2                               stringsAsFactors = FALSE);
```

And we are going to save the weather data columns to vectors.

```
1  date = lansing2016weather[, "date"];  
2  eventData = lansing2016weather[, "eventData"];  
3  avgTemp = lansing2016weather[, "avgTemp"];  
4  tempDept = lansing2016weather[, "tempDept"];  
5  precipitation = lansing2016weather[, "precipitation"];  
6  humidity = lansing2016weather[, "humidity"];  
7  barometer = lansing2016weather[, "barometer"];  
8  dewPoint = lansing2016weather[, "dewPoint"];  
9  avgwind = lansing2016weather[, "avgwind"];  
10 maxwind = lansing2016weather[, "maxwind"];  
11 windDirection = lansing2016weather[, "windDirection"];  
12 sunrise = lansing2016weather[, "sunrise"];  
13 sunset = lansing2016weather[, "sunset"];
```

### 4 - which() to subset vectors

As we saw in the last lesson, we can find out which days had an average temperature greater than 80 using the **which()** function:

```
1  daysOver80 = which(avgTemp > 80);
```

**which()** produces a vector of index values (**163, 204, 224, 225, 250**) and this vector is saved to the variable

**daysOver80.** The five vector values indicate that there were **5 days** with average temperatures greater than 80 (the 163rd day, the 204th day...).

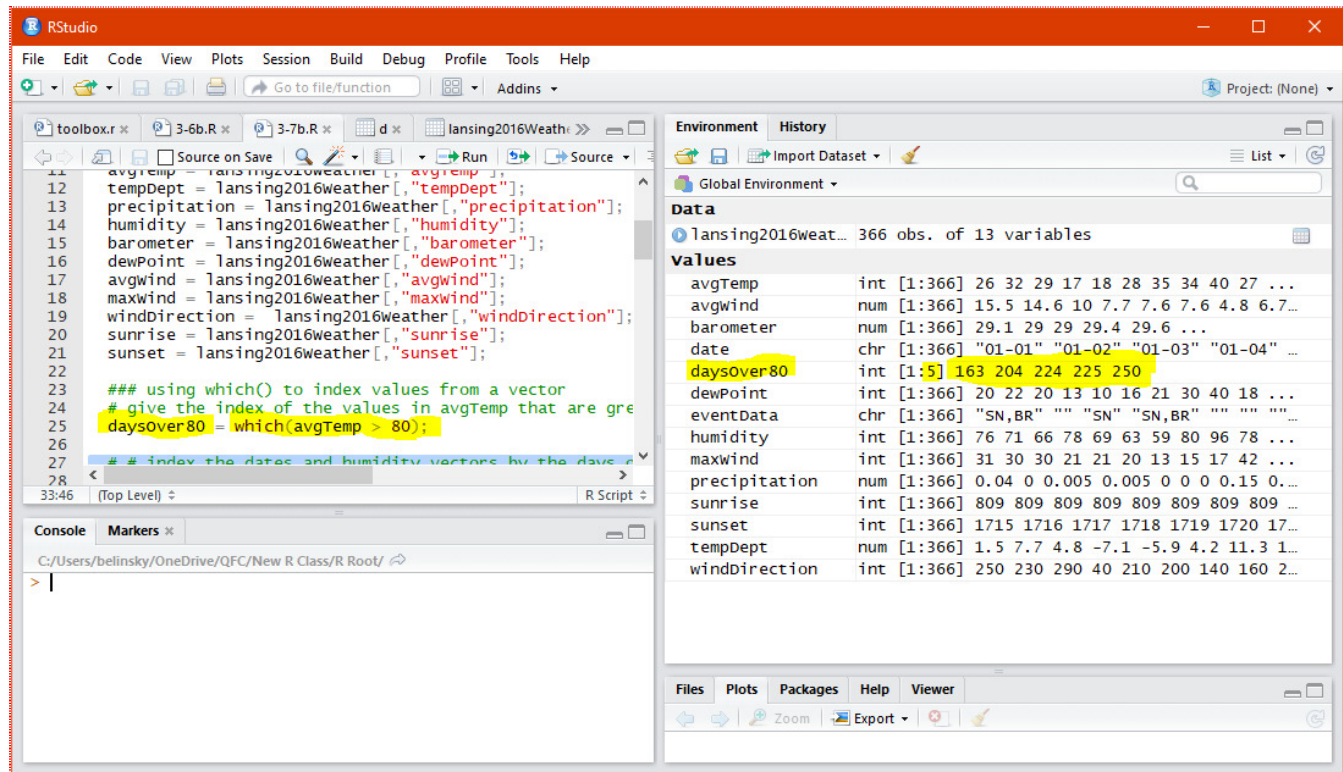


Fig 1: Using **which()** to subset the vector **avgTemp**.

## 4.1 - Indexing other vectors

**daysOver80** is a vector of five index values with:

- **daysOver80[1]** = 163
- **daysOver80[2]** = 204
- **daysOver80[3]** = 224
- **daysOver80[4]** = 225
- **daysOver80[5]** = 250

We can use the vector **daysOver80** to index other vectors and find out more information like:

What are the actual dates where the temperatures averaged over 80 degrees?

```
1 datesOver80 = dates[daysOver80];
```

*answer: June 11, July 22, August 11, August 12, and September 6*

Or, what was the humidity on the days the temperature averaged over 80 degrees?

```
1 humidityOver80 = humidity[daysOver80];
```

*answer: 54%, 65%, 66%, 82%, and 70%*

## 4.2 - Simple stats on subset vector

We could take this one step further and find the **min()** and **max()** humidity for the days the temperatures averaged over 80 degrees.

```
1 minHumOver80 = min(humidity[daysOver80]);
```

```
2 | maxHumOver80 = max(humidity[daysOver80]);
```

*answers: min=54%, max=82%*

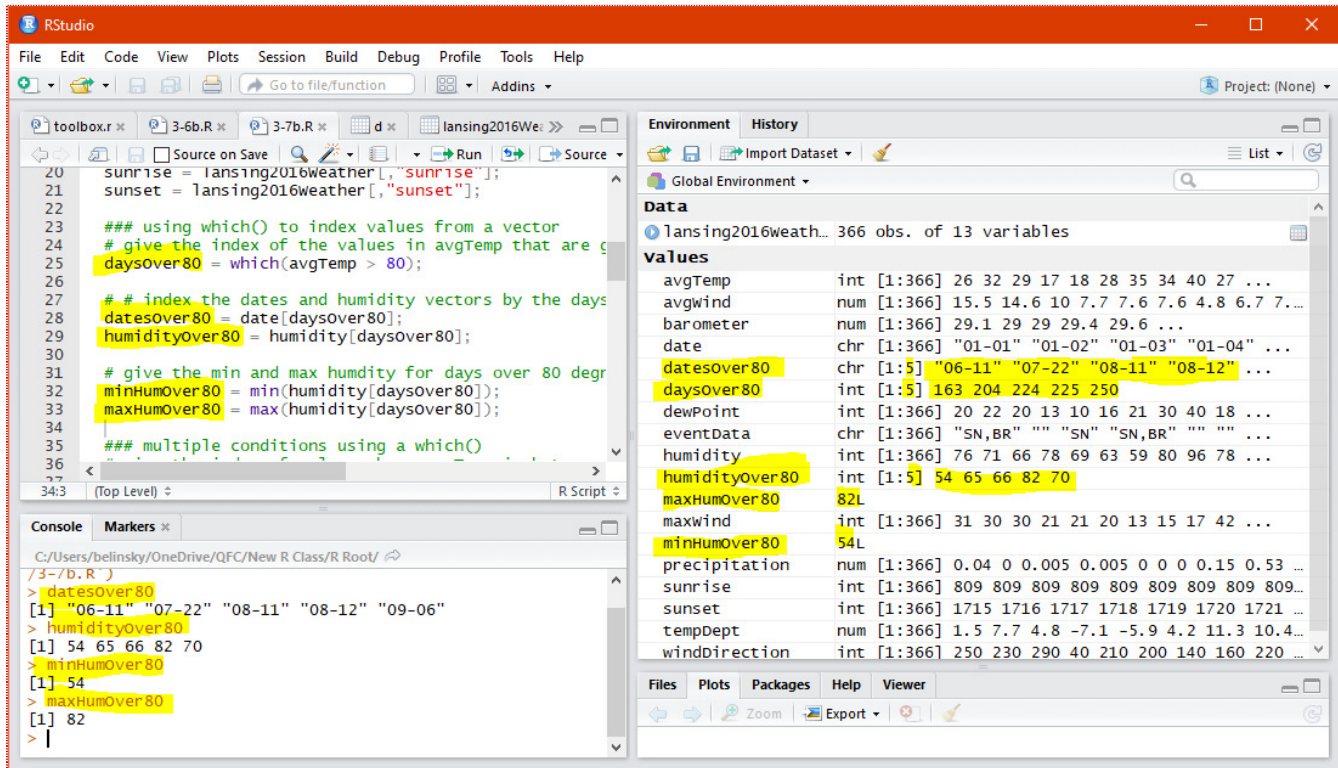


Fig 2: Use the vector of index values (**daysOver80**) to subset other vectors (**dates** and **humidity**).

### 4.3 - Creating index vector with multiple conditions (& and |)

Let's find the days that averaged temperatures in the **60s** (so, between **60** and **70**).

For *a single temperature value*, we could use an **if()** with the **&&** operator:

```
1 | if(avgTemp > 60 && avgTemp < 70)
```

But we have a vector with multiple values. We can make a similar conditional statement using **which()**, except we use a single **&** instead of the double **&&**.

**Extension: Difference between & and &&**

To find temperatures between **60** and **70** using **which()** we look for **avgTemp** greater than **60** and less than **70**. As with the **if()** conditional statement, we need to explicitly give both conditions.

```
1 | tempsIn60s = which(avgTemp > 60 & avgTemp < 70);
```

**tempsIn60s** has 56 values (**Fig 3**), so there were 56 days that averaged 60 to 70 degrees. The earliest day was day **69** (early March) and the latest was day **306** (early November).

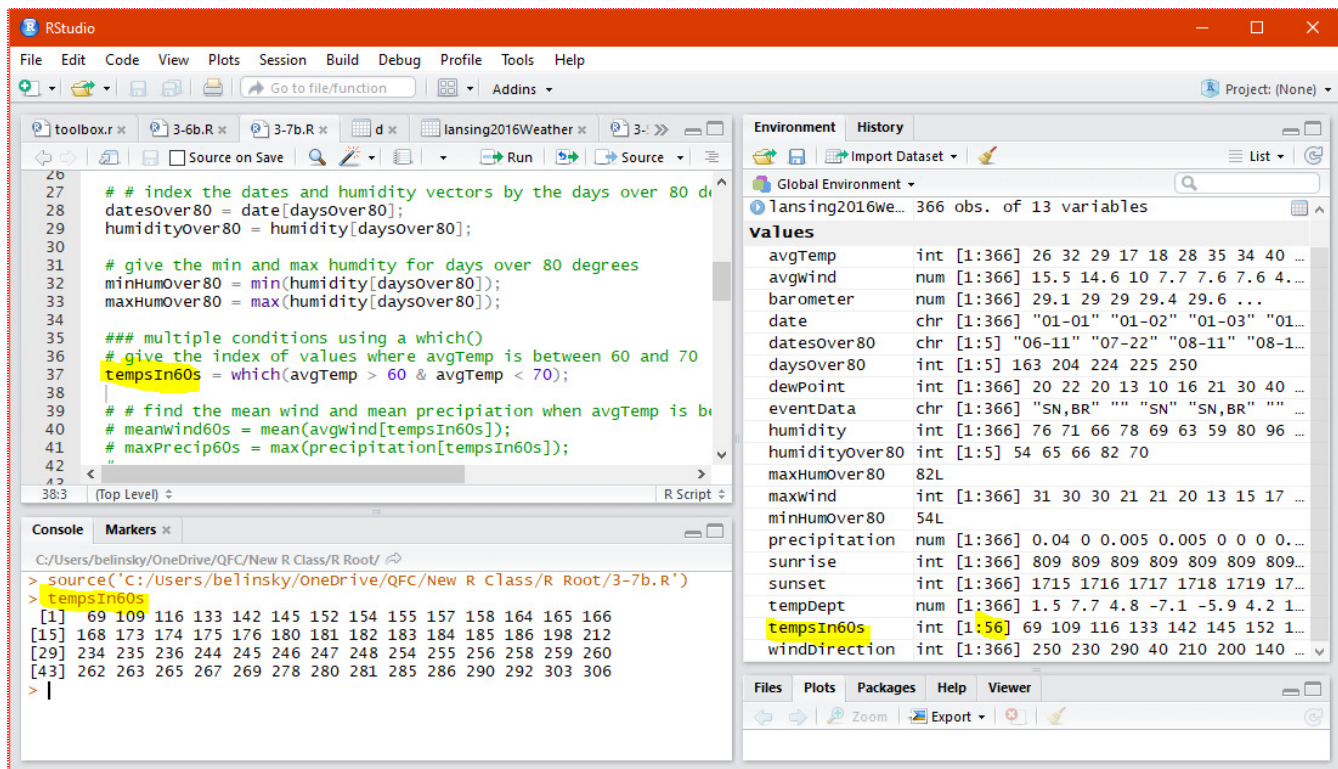


Fig 3: Using the `&` operator in `which()`.

Now we can use the vector ***tempsIn60s*** to find the average winds for the **56** days with temperature between **60** and **70**:

```
1 | meanWind60s = mean(avgWind[tempsIn60s]);
```

or the maximum precipitation for the **56** days with high temperature between **60** and **70**:

```
1 | maxPrecip60s = max(precipitation[tempsIn60s]);
```

We can see that (Fig 4):

- there were **56** days with ***avgTemp*** values in the **60s**
- the average wind speed **for those 56 days** was about **7.3mph**
- and the maximum precipitation for those **56** days was **1.01 inches**.





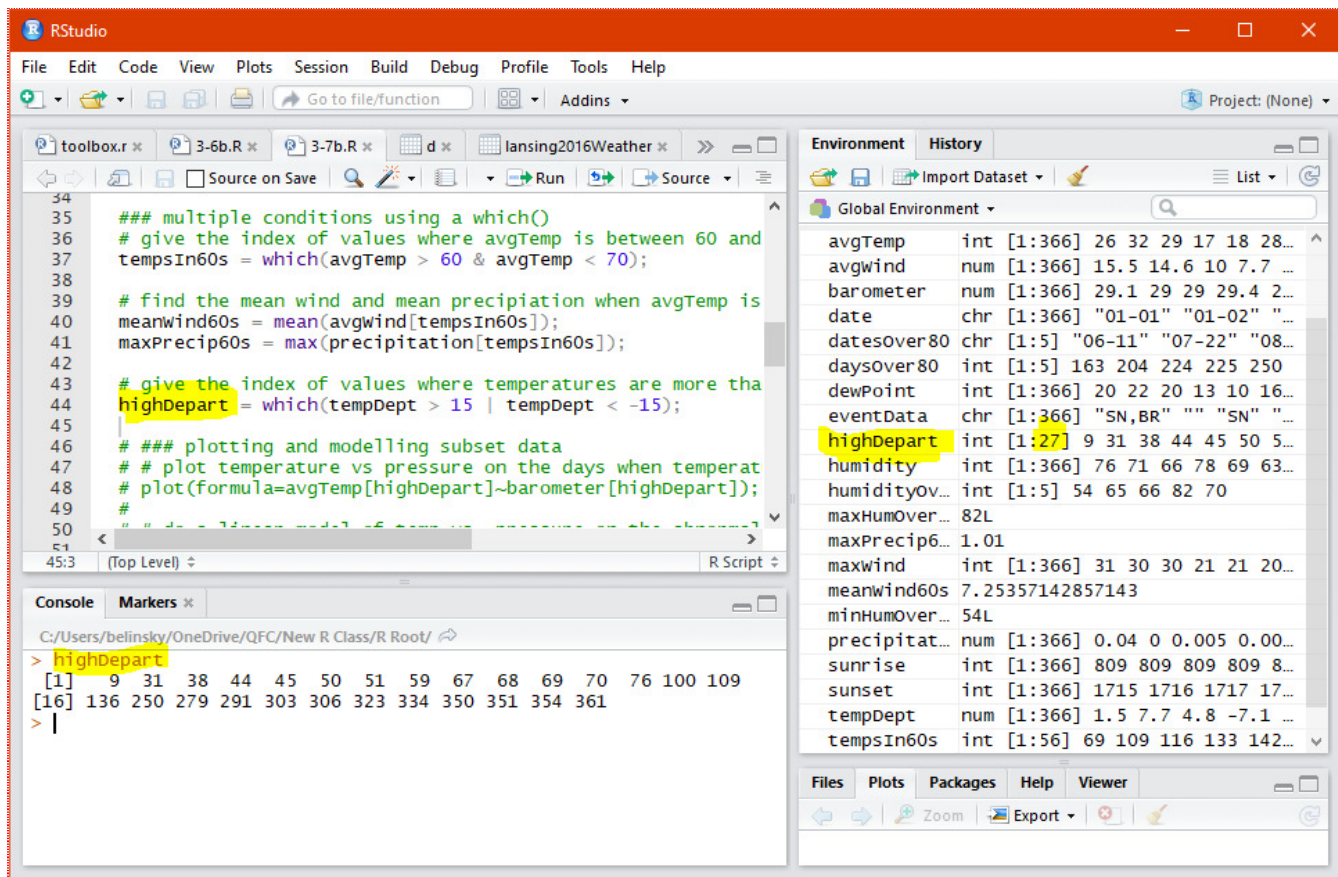


Fig 5: Subsetting a vector using `which()` and the `or (|)` conditional operator.

## 5.1 - Set up scatterplot and linear model

Next we make a scatterplot of **avgTemp** vs **barometer** on days with high departure temperatures (Fig 6):

```
1 | plot(formula=avgTemp[highDepart]~barometer[highDepart]);
```

And make a linear model that indexes temperature and pressure on the days with more than a 15 degree deviation from normal (Fig 6):

```
1 | model = lm(formula=avgTemp[highDepart]~barometer[highDepart]);
```

Then we add the regression line from the model to the scatterplot (Fig 6):

```
1 | abline(model, col="blue");
```

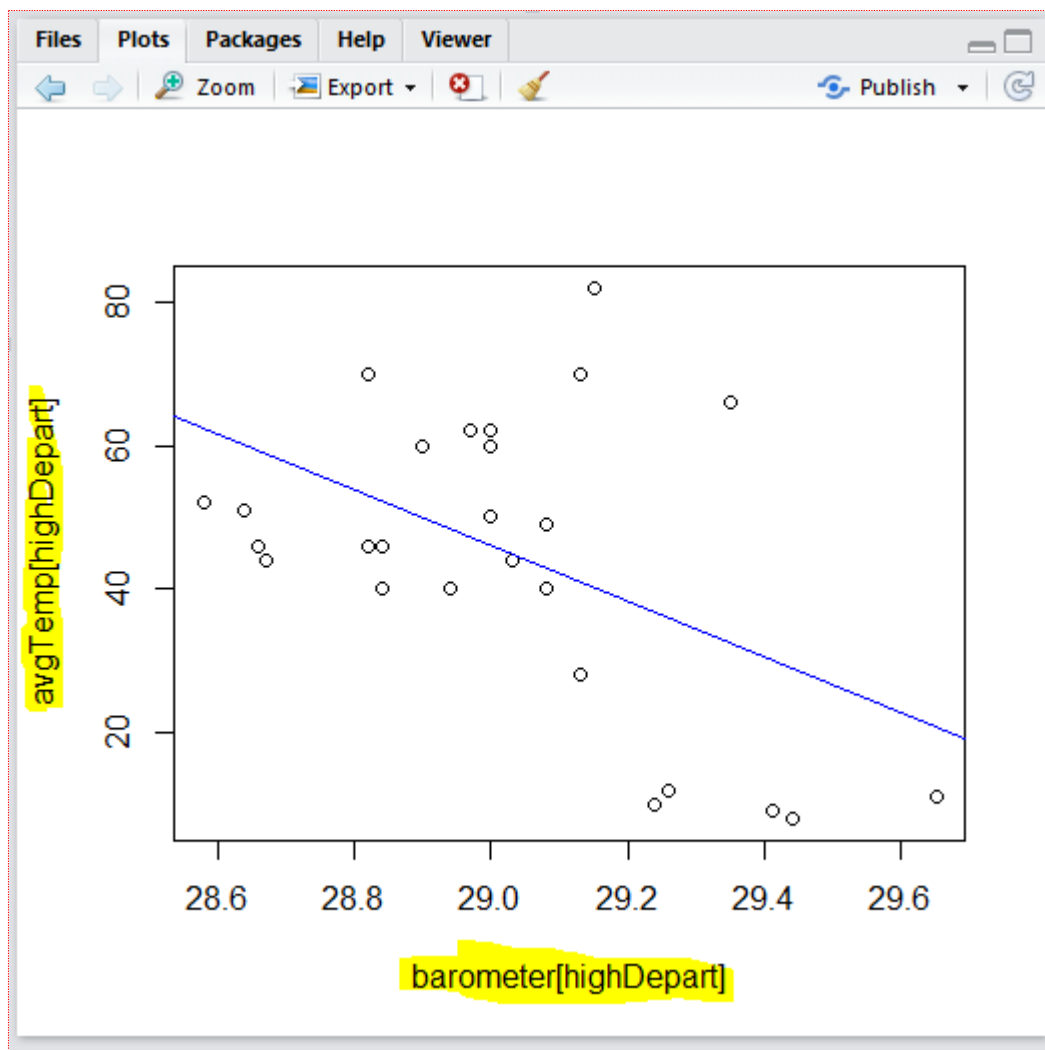


Fig 6: Scatterplot of **avgTemp** vs **barometer** on days with abnormal temperatures -- with regression line added

The scatterplot and regression suggest a relationship between **avgTemp** and **barometer** when temperatures deviates by more than 15 degrees from normal.

## 5.2 - Summarizing the linear model

We can print the linear model summary to the Console Window

```
1 | print(summary(model));
```

```

Console Markers x
C:/Users/belinsky/OneDrive/QFC/New R Class/R Root/
> source("C:/Users/belinsky/OneDrive/QFC/New R Class/R Root/3-7b.R")

Call:
lm(formula = avgTemp[highDepart] ~ barometer[highDepart])

Residuals:
    Min       1Q   Median       3Q      Max
-26.75 -12.63  -6.24   14.37   41.76

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)    1168.81     401.68   2.910  0.00749 **
barometer[highDepart] -38.72      13.84  -2.797  0.00977 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18.31 on 25 degrees of freedom
Multiple R-squared:  0.2384,    Adjusted R-squared:  0.2079
F-statistic: 7.826 on 1 and 25 DF, p-value: 0.009769

> |

```

Fig 7: Summary of the linear model that suggests a relationship between **avgTemp** and **barometer**.

### 5.3 - A second linear model

Wind directions plays a big role in weather conditions so we will repeat the analysis with southerly winds (**windDirection** between **90 and 270** degrees). Note: **0** degrees is due north, **180** degrees is due south.

```

1 southwinds = which(windDirection > 90 & windDirection < 270);
2 plot(formula=avgTemp[southwinds]~barometer[southwinds]);
3 model2 = lm(formula=avgTemp[southwinds]~barometer[southwinds]);
4 abline(model2, col="blue");    # add the regression line to the plot
5 print(summary(model2));        # summary shows a relationship

```



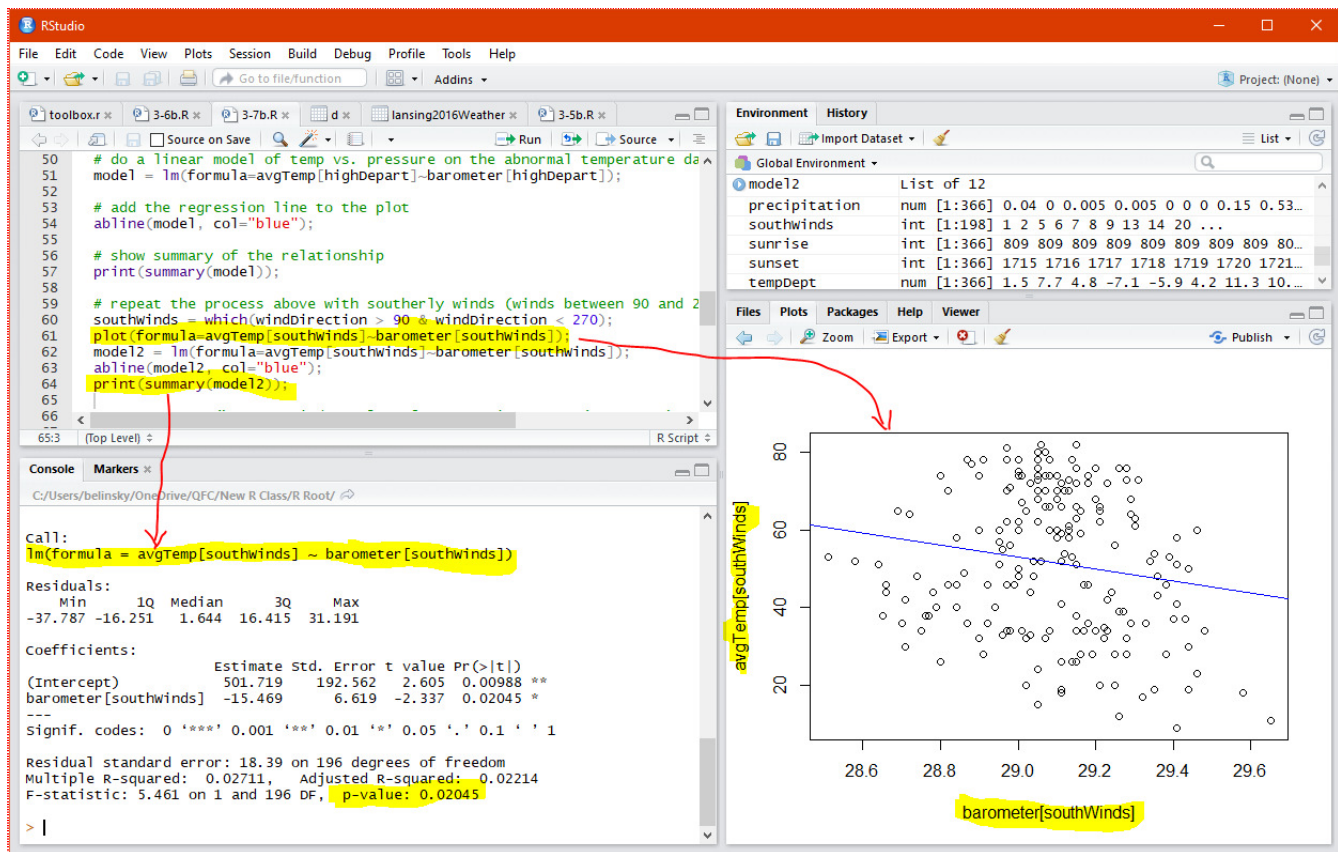


Fig 8: Summary of linear model showing relationship between **avgTemp** and **barometer** on days with southerly winds

## 6 - Finding patterns in string value with grep()

**which()** works well for subsetting vectors that have numeric values but **which()** is not equipped to deal flexibly with string values. As we saw in last lesson, **eventData** has values like: **SN**, **FG**, **HZ**, indicating **snow**, **fog**, and **haze** on that day.

We can use **grep()** to find all **eventData** values that contain **SN** (snow) (Fig 9):

```
1 dayswithSnow = grep("SN", eventData);
```

or use **grep()** find all **eventData** values that contain **RN** (rain) (Fig 9):

```
1 dayswithRain = grep("RN", eventData);
```

### 6.1 - grep() with the | condition

And we can use the **or ( / )** operator to find all values in **eventData** that contain either **SN** or **RN** (Fig 9):

```
1 dayswithPrecip = grep("RN|SN", eventData);
```

There are **124** days with only rain, **65** with only snow, and **179** days with rain **or** snow.

**124 + 65 = 189**, so there are **10** days (**189-179**) that had **both** rain **and** snow.

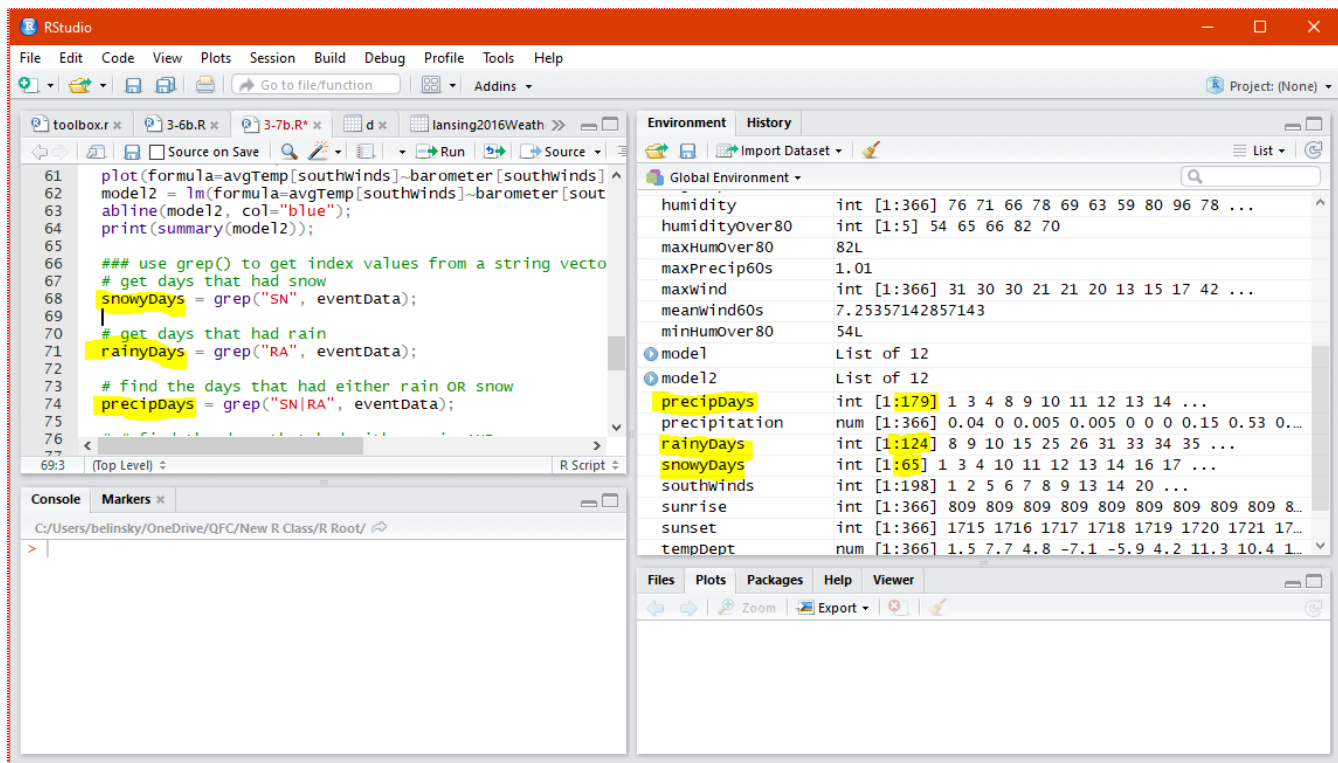


Fig 9: Using the or ( | ) condition within a `grep()`.

We can return days with rain or snow using `grep()` (the | operator) but we cannot use `grep()` to return days with **both** rain **and** snow (the & operator)

```
1 daysWithRainAndSnow1 = grep("RN&SN", eventData); # this does not work!
```

The reason for this requires a lengthy explanation of how `grep()` (and regular expressions, in general) are executed.

## 6.2 - An alternative to & when using grep()

We have:

- **daysWithSnow**: the indexes of all days with snow, and
- **daysWithRain**: the indexes of all days with rain.

What we want is the *index values that are in both daysWithRain and daysWithSnow*.

Luckily, there is a function for that called `intersect()`. `intersect()` produces a vector of values that occur in both vectors.

```
1 daysWithRainAndSnow2 = intersect(daysWithRain, daysWithSnow);
```

**daysWithRainAndSnow2** has 10 values, representing the 10 values in **eventData** that had both **RN** (rain) and **SN** (snow).

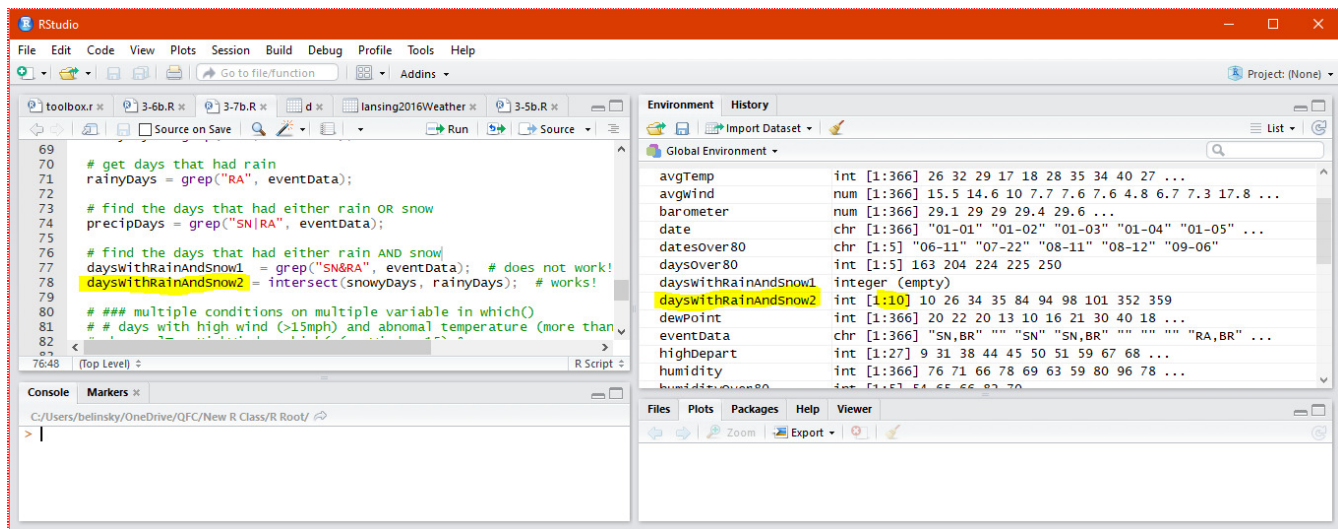


Fig 10: Using `intersect()` to find values that occur in two vectors.

Extension: `union()` -- sister function to `intersect()`

## 7 - Multiple conditions on multiple variables

Similar to `if()`, we can use `which()` to look at *multiple conditions on multiple variables*. The trick is to make sure you have the parentheses correct, because parentheses determine the order of operations.

`abnormalTempHighWind` gives the index of values where there were strong winds (greater than **15mph**) and the temperature deviated more than **10** degrees from average.

```

1 abnormalTempHighWind = which( (avgWind > 15) &
2                               (tempDept < -10 | tempDept > 10 ) );

```

`chillyLightRain` gives the index for temperature values between **40 and 50** degrees where there was a light precipitation (between **0.1 and 0.2** inches)

```

1 chillyLightRain = which( (avgTemp > 40 & avgTemp < 50) &
2                           (precipitation > 0.1 & precipitation < 0.2 ) );

```

We see there are:

- 9 windy days with temperatures deviating more than 10 degrees from average and
- 3 days that were chilly with light rain.

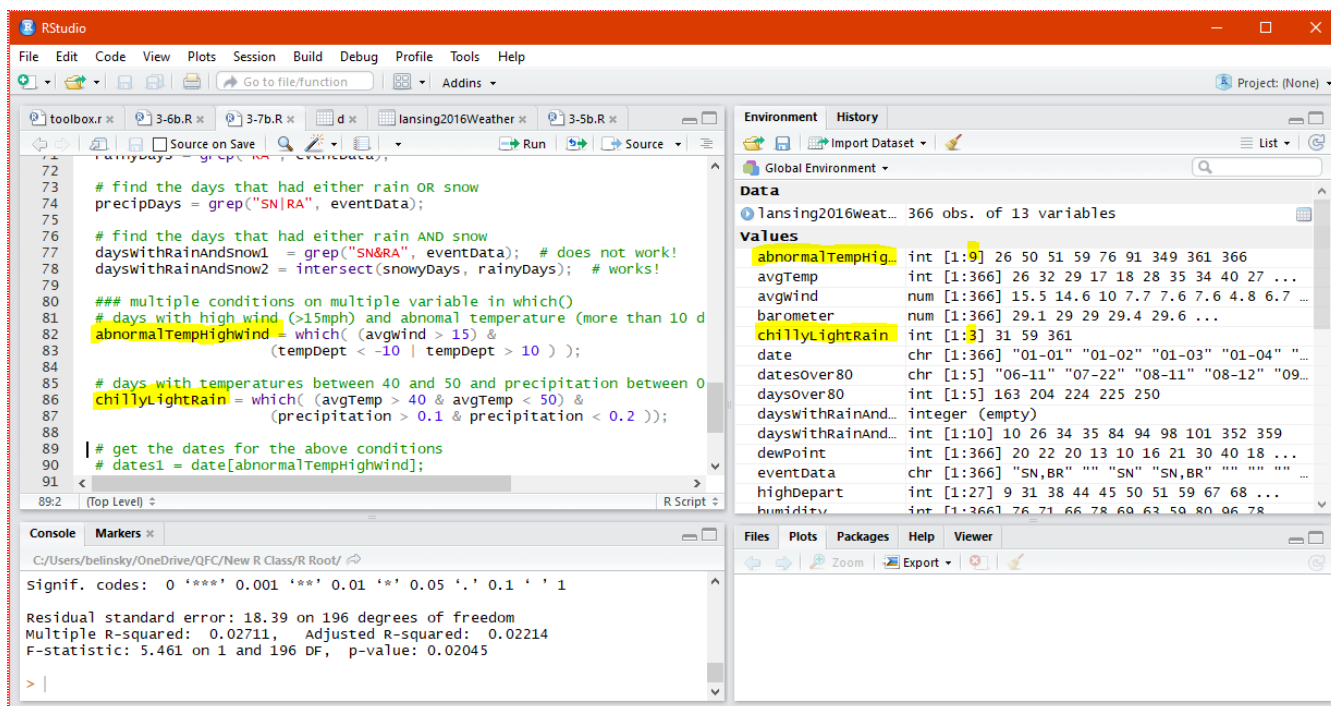


Fig 11: Using **which()** with multiple conditions on multiple variables

## 7.1 - Using the index vector to subset another vector

Let's get the actual dates for the conditions above.

We can use the subset vectors to find the dates for **abnormalTempHighWind** and **chillyLightRain**:

```

1 dates1 = date[abnormalTempHighWind];
2 dates2 = date[chillyLightRain];

```

And we will print this information to the Console Window. Note: **cat()** is easier to use than **print()** for printing out string-only content.

```

1 cat("Abnormal temps and high winds:\n");
2 print(dates1);
3 cat("\nChilly days and light rain:\n");
4 print(dates2);

```



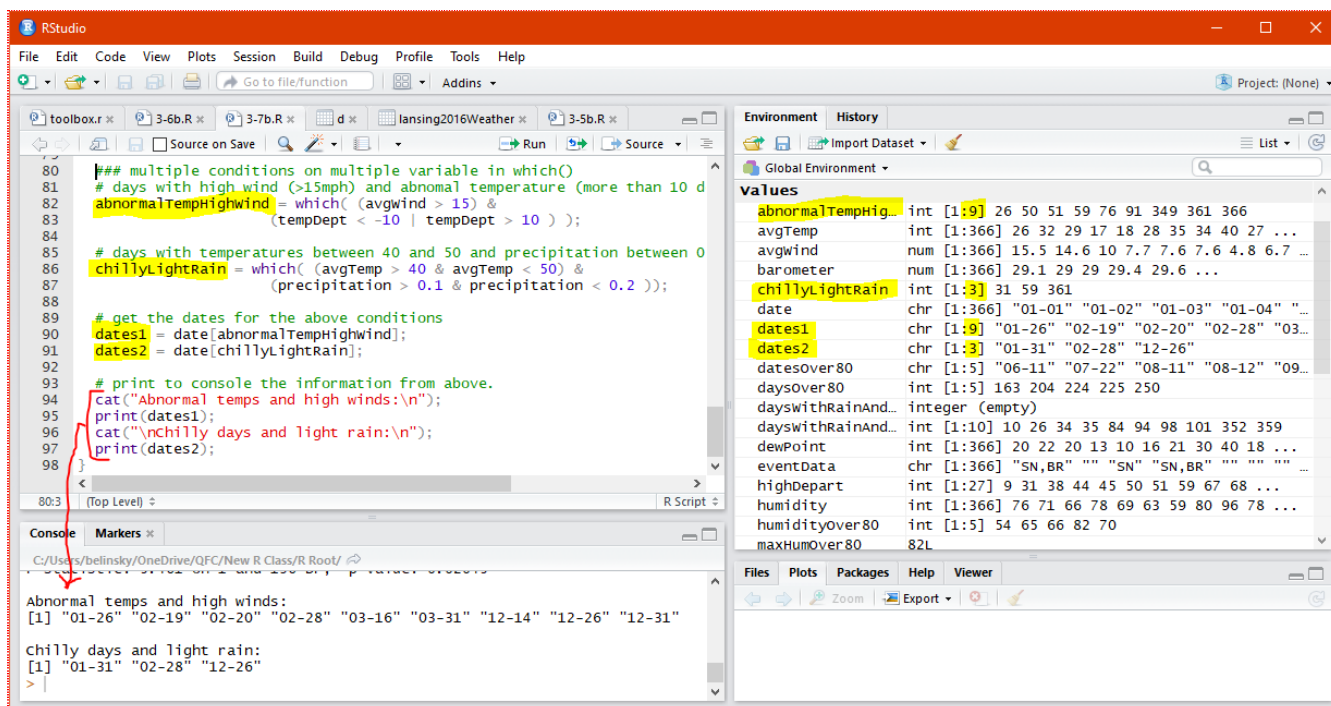


Fig 12: Printing out the results of subsetting vectors

## 8 - Application

Use the data from **LansingNOAA2016Formatted.csv** for this application.

- Using **grep()**, create a vector with indices for every day in the months of February and March.  
hint: all December values are in the form **12-##**, all November values are in the form **11-##...**
- Find the number of days the maximum wind speed was high (greater than 40) and for these days give:
  - the date
  - the amount of precipitation
  - the maximum wind speed
- Print out the dates that have
  - northerly winds (less than 90 degrees or greater than 270 degrees)
  - northerly winds and precipitation greater than 0.5 inches
- Does **hoursOfSun** (from last lesson's application) always correlate with temperature?
  - Make a scatterplot of **avgTemp** vs **hoursOfSun** for every two months (6 in all)
  - Create a linear models of **avgTemp** vs **hoursOfSun** for every two months (6 in all)
  - Add the regression lines to the scatterplots

## 9 - Extension: Differences between & and &&

Note: everything in this extension is also true for the differences between | and ||.

If you are comparing single values then **&** and **&&** will always function exactly the same so:

```

1 testNum = 4
2

```

```

3   if(testNum > 3 && testNum < 7) # evaluates to TRUE
4   {
5       # do something here
6   }
7
8   if(testNum > 3 & testNum < 7)  # also evaluates to TRUE
9   {
10      # do something here
11  }

```

Most R programmers will use **&** in these circumstances. For this class, I chose to use **&&** because *in most other programming languages you need to use &&* when you are comparing single values. However, you could go back to lessons 1-9 and 1-10 and replace every **&&** with **&** and nothing will functionally change.

When dealing with vectors in R, **&** and **&&** will function differently. **&&** checks to see if *every value in the vector meets the condition* and produces one **TRUE** or **FALSE** statement whereas **&** *checks each value individually* and produces a **TRUE/FALSE** statement for each value.

Let's look at an example using the vector **ages = c(2,7,3,9,6,3,5)**:

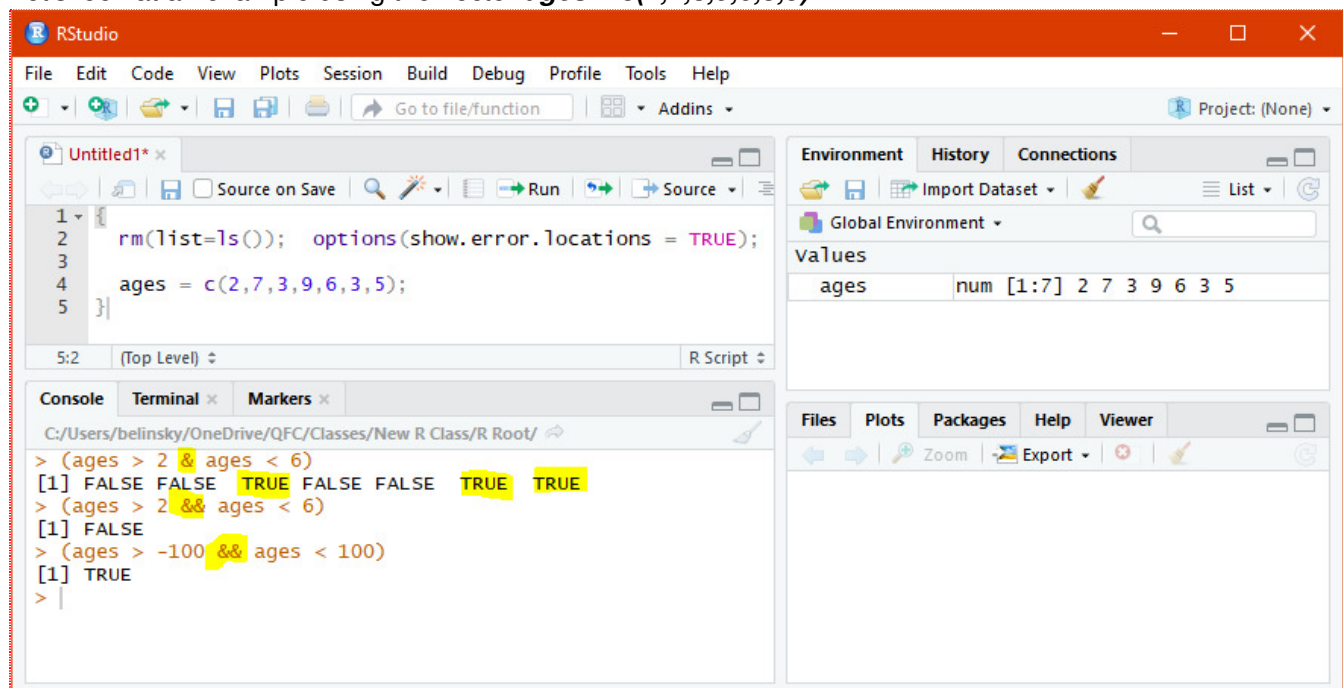


Fig 13: Comparing **&** to **&&** in conditional statements with vectors.

In the Console Window we see that **(age > 2 & age < 6)** produces seven values, representing each value of the **ages** vector. Since the 3<sup>rd</sup>, 6<sup>th</sup>, and 7<sup>th</sup> values in **ages** are between 2 and 6, the 3<sup>rd</sup>, 6<sup>th</sup>, and 7<sup>th</sup> values are **TRUE**, the rest are **FALSE**.

**(age > 2 && age < 6)** produces one **FALSE** value because all the values in **ages** *do not meet the condition of being between 2 and 6*.

**(age > -100 && age < 100)** produces one **TRUE** value because all the values in **ages** *do meet the condition of being between -100 and 100*.

Note: In most other programming language, **&** functions differently -- it is a *bitwise operator*, which is an advanced topic.

## 10 - Extension: union(), the sister to intersect()

**intersect()** on two vectors returns the values that are represented in *both vectors*,  
**union()** on two vectors returns the values that are represented in *either vector*.

So, if:

```
1 aVector = c(3, 4, 5, 6, 7);
2 bVector = c(6, 7, 8, 9, 10);
```

then:

```
1 intersect(aVector, bVector);
```

returns a vector with the values **6** and **7**, because **6** and **7** occurred in both **aVector** and **bVector**

and:

```
1 union(aVector, bVector);
```

returns a vector with the values **3,4,5,6,7,8,9,10** -- or all the values that occur in either **aVector** or **bVector**

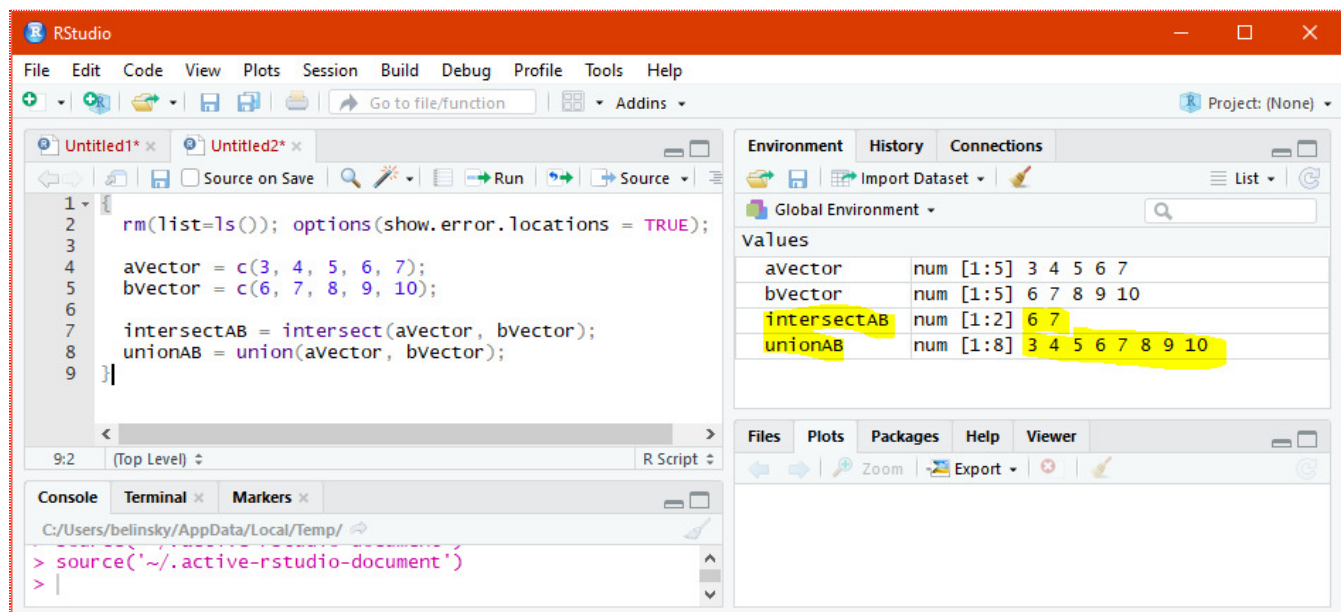


Fig 14: The intersection and union of two vectors.